

Importing the necessary libraries

```
In [1]: import tensorflow as tf
import tensorflow.keras as keras
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import cv2 as cv
```

Splitting the data into train and test data

- 60k -> training data
- 10k -> testing data

```
In [2]: (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
```

Normalizing the data

28 * 28 grid being converted to 784 single column

```
In [3]: x_train = keras.utils.normalize(x_train, axis=1)
x_test = keras.utils.normalize(x_test, axis=1)
```

Making the deep learning neural network with 2 hidden layers

```
In [4]: model = keras.models.Sequential()
model.add(keras.layers.Flatten(input_shape = (28, 28)))
```

Making of 2 hidden layers using relu as the activation function

```
In [5]: model.add(keras.layers.Dense(units=128, activation=tf.nn.relu))
model.add(keras.layers.Dense(units=128, activation=tf.nn.relu))
```

The output layer of 10 neurons (0-9)

```
In [6]: model.add(keras.layers.Dense(units=10, activation=tf.nn.softmax))
```

Compiling the model with adam optimizer

```
In [7]: model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

Fitting the data

```
In [8]: model.fit(x_train, y_train, epochs=5)
```

```
Epoch 1/5
1875/1875 [=====] - 4s 2ms/step - loss: 0.2615 - accuracy: 0.9238
Epoch 2/5
1875/1875 [=====] - 3s 2ms/step - loss: 0.1061 - accuracy: 0.9673
Epoch 3/5
1875/1875 [=====] - 3s 2ms/step - loss: 0.0745 - accuracy: 0.9769
Epoch 4/5
1875/1875 [=====] - 3s 2ms/step - loss: 0.0542 - accuracy: 0.9823
Epoch 5/5
1875/1875 [=====] - 3s 2ms/step - loss: 0.0406 - accuracy: 0.9868
<keras.callbacks.History at 0x1b8543618e0>
```

Calculating the loss and the accuracy of the model on the testing dataset


```
In [9]: loss, accuracy = model.evaluate(x_test, y_test)
print("Loss: ", loss)
print("Accuracy: ", accuracy)
```

```
313/313 [=====] - 0s 1ms/step - loss: 0.0935 - accuracy: 0.9739
Loss: 0.09354386478662491
Accuracy: 0.9739000201225281
```


Running the model on the hand-written images (made on paint)

```
In [13]: for x in range(1,7):
img = cv.imread(f'{x}.png')[ :, :,0]
img = np.invert(np.array([img]))
prediction = model.predict(img)
print(f'The predicted result is: {np.argmax(prediction)}')
plt.imshow(img[0], cmap=plt.cm.binary)
plt.show()
```


The predicted result is: 1




The predicted result is: 2




The predicted result is: 3




The predicted result is: 5



The predicted result is: 7



The predicted result is: 8



In []: